



PROG8080

Data Types and CREATE TABLE in SQL Server 2014

Glenn Paulley
Fall 2015

PART 1: SQL DATA TYPES

Data Types

- Each column in a table has a data type
- The data type determines the type and amount of data you can store in a column
 - Different types have different semantics and support different sets of operations

Character Types

- CHAR(*n*)
- VARCHAR(*n*)
- TEXT
- NCHAR(*n*)
- NVARCHAR(*n*)
- nTEXT
- Fixed-length, up to 8K characters
- Variable length, up to 8K characters
- Variable length, up to 2GB characters
- Fixed length Unicode data, up to 4000 characters
- Variable length Unicode data, up to 2GB
- Variable length Unicode data, up to $2^{30} - 1$ characters



Binary (bit string) Types

- `BINARY(n)`
- `VARBINARY(n)`
- `IMAGE`
- `BIT`
- Fixed-length, up to 8K bytes
- Variable length, up to 8K bytes
- Variable length, up to 2GB bytes
- Supported values are 0,1, or NULL

Exact Numeric Types

- INT
 - Whole numbers $\pm 2^{31}$
 - 32-bit signed values
- BIGINT
 - Whole numbers $\pm 2^{63}$
 - 64-bit signed values
- SMALLINT
 - Whole numbers $\pm 32K$
 - 16-bit signed values
- TINYINT
 - Positive whole numbers 0-255



DECIMAL and NUMERIC types

- DECIMAL and NUMERIC are synonyms
- DECIMAL is an exact numeric type that supports numeric values with fractional component
- Declaration specifies *precision* (between 1 and 38) and an optional *scale*
 - DECIMAL(7,2)
 - 7 significant digits (precision), 2 of these digits after the decimal point (scale)
 - DECIMAL(10)
 - 10 significant digits, no fractional component
- Range: +/- 10³⁸
- Storage: between 5 and 17 bytes
- Default: if a value is declared DECIMAL, assumed to be DECIMAL(18)

Approximate Numeric Types

- Approximate numeric types store floating-point values; specification uses scientific notation (mantissa and exponent)
 - Not all numbers can be represented exactly (mantissa is in base-2)
 - Should never be used for primary keys
- REAL
 - Storage: 4 bytes
 - Range: $-3.40E + 38$ to $-1.18E - 38$, 0 and $1.18E - 38$ to $3.40E + 38$
- FLOAT(n)
 - Value of n is between 1 and 53; number of bits used to store the mantissa of the number
 - Storage: 4 bytes if $n \leq 24$; 8 bytes if $n > 24$
- DOUBLE PRECISION
 - Synonym for FLOAT(53)



Date/Time types in SQL Server 2014

- DATE
- TIME
- DATETIME
- SMALLDATETIME
- DATETIME2
- DATETIMEOFFSET
- Note: No TIMESTAMP type in this list (from the SQL Standard)
 - The TIMESTAMP type means something else in SQL Server!



DATE type

- Range: 1 January 0001 AD through 31 December 9999
 - Cannot be used for dates BC
- Date only, no time portion
- Default string format: 'YYYY-MM-DD'
- Storage size: 3 bytes
- Default value: 1 January 1900



TIME type

- Range: 00:00:00.00000000 through 23:59:59.99999999
 - A TIME type can be declared with millisecond precision
 - Default is TIME(7), 100ns precision
- Time only, no date portion
- Default string format: 'hh:mm:ss[.nnnnnnnn]'
- Storage size: between 3 and 5 bytes, depending on precision
- Default value: 00:00:00



DATETIME type

- Range: January 1, 1753, through December 31, 9999
 - Time portion is between 00:00:00 and 23:59:59.997
 - Approximately 1/300 second precision with the time portion
- Character length: 19 minimum to 23 characters maximum
- Default string format: 'YYYY-MM-DD hh:mm:ss[.nnnnnnnn]'
- Storage size: 8 bytes
- Default value: 1900-01-01 00:00:00.000
- SET DATEFORMAT option controls interpretation of DATETIME values when specified as string literals



SMALLDATETIME type

- Range: January 1, 1900, through June 6, 2079
 - Time portion is between 00:00:00 and 23:59:59
 - Approximately 1 minute precision with the time portion
- Character length: 19 characters maximum
- Default string format: 'YYYY-MM-DD hh:mm:ss'
- Storage size: 4 bytes
- Default value: 1900-01-01 00:00:00

DATETIME2 type

- SQL Server 2014's implementation of the TIMESTAMP type from the ISO SQL Standard
 - Now recommended for all applications instead of DATETIME
- Range: January 1, 0001, through December 31, 9999
 - Cannot be used for dates BC
 - Time portion is between 00:00:00 and 23:59:59.9999999
 - DATETIME2 supports specification of the time precision, up to 100ns
 - DATETIME2(7) is the default (100ns precision)
- Character length: 19 minimum to 27 characters maximum
- Default string format: 'YYYY-MM-DD hh:mm:ss[.fractional seconds]'
- Storage size: between 6 and 8 bytes depending on time precision
- Default value: 1900-01-01 00:00:00.0000000



DATETIMEOFFSET type

- SQL Server 2014's implementation of the TIMESTAMP WITH TIME ZONE type from the ISO SQL Standard
 - Adds a time zone component to a DATETIME2 type
- Range: January 1, 0001, through December 31, 9999
 - Cannot be used for dates BC
 - Time portion is between 00:00:00 and 23:59:59.9999999
 - DATETIMEOFFSET supports specification of the time precision, up to 100ns
 - DATETIME2(7) is the default (100ns precision)
 - Time zones: -14h through +14h
- Character length: 26 minimum to 34 characters maximum
- Default string format: 'YYYY-MM-DD hh:mm:ss.nnnnnnnn {+|-}hh:mm
- Storage size: 10 bytes
- Default value: 1900-01-01 00:00:00.00000000 +00.00



Other (vendor-specific) Types

- MONEY
 - Money amounts to the trillions with 4 decimal places accuracy
- ROWVERSION
 - (formerly known as TIMESTAMP)
 - A database-wide unique row number
 - TIMESTAMP is still supported as a row version type

User Defined Data Types

- It is possible to define user-defined types (UDTs)
 - UDTs are existing types with specific properties
- UDTs are an advanced technique used to ensure consistency across tables



PART 2: CREATE TABLE

CREATE TABLE Syntax

- Basic CREATE TABLE syntax and format:

```
CREATE TABLE tableName  
( column1 TYPE [qualifiers],  
  column2 TYPE,  
  column3 TYPE [qualifiers],  
  ...  
  [, table-level constraints]  
)
```

- It is customary, but not essential, to put primary key columns at the beginning of a table

CREATE TABLE Example 1

```
CREATE TABLE product
(
  product_id VARCHAR(11),
  description VARCHAR(75),
  vendor_id VARCHAR(4),
  vendor_part_number VARCHAR(20),
  price NUMERIC(7,2),
  reorder_threshold INT,
  product_category_code CHAR(1)
)
```



NULL and NOT NULL

- Each column definition can include:
 - NULL to indicate that NULL values are permitted
 - This is the default and is therefore seldom seen
 - NOT NULL to indicate that NULL values are prohibited

...

vendor_id VARCHAR(4) NOT NULL,

...

CREATE TABLE

- All database systems have various optional clauses on CREATE TABLE that make these statements non-portable across various database management systems
- CREATE TABLE can contain vendor-specific syntax for:
 - Column and/or table-level encryption
 - Column compression
 - FOREIGN KEY constraints and various options
 - Table partitioning specification
 - Indexing options
 - Use of system-generated values for primary keys
 - SQL Server 2014 supports both IDENTITY values and SEQUENCE types



PART 3: CREATING A TABLE FROM A SELECT STATEMENT



CREATE TABLE

- In SQL Server, CREATE TABLE can only be used to create a table from scratch
 - Populating a table afterwards can be done using an INSERT INTO <table> FROM SELECT ...
- This differs from ORACLE, for example, where one can specify
CREATE TABLE myCourses AS
SELECT * FROM Course
- To do this in SQL Server, one uses a SELECT statement with an INTO clause



SELECT INTO

- Syntax:

SELECT <select list>

INTO <table>

FROM ...

WHERE ...

GROUP BY ...

HAVING ...

- The columns in <table> have the same names and data types as the attributes in the query's SELECT list
- Rows are populated from the result of the SELECT statement



SELECT INTO - limitations

- Cannot be used to create a partitioned table
- Indexes, constraints, and triggers defined in the source table are not transferred to the new table, nor can they be specified in the SELECT...INTO statement.
 - If these objects are required, you can create them after executing the SELECT...INTO statement.
- Specifying an ORDER BY clause does not guarantee the rows are inserted in the specified order.
- When a computed column is included in the select list, the corresponding column in the new table is not a computed column. The values in the new column are the values that were computed at the time SELECT...INTO was executed.



PART 4: DROP TABLE AND IF EXISTS



DROP TABLE

- To completely remove a table from a database, use DROP TABLE
 - DROP TABLE table

IF EXISTS

- Use code like this in a script to drop a table if it already exists:

```
IF EXISTS ( SELECT NAME FROM sysobjects
            WHERE name = 'mytable' )
DROP TABLE mytable;
```

or conversely

```
IF OBJECT_ID('dbo.your_table', 'U') IS NOT NULL
DROP TABLE dbo.your_table;
```

